
Korean 4-Ball Billiards: A Continuous, Deterministic, Sparse-Reward Benchmark Solved by Inference-Time Search

Doyeol Oh
KAIST
ohdoyeol@kaist.ac.kr

Byungmo Kang
KAIST
kbm03@kaist.ac.kr

Seojun Park
KAIST
tjwnsl8@kaist.ac.kr

Abstract

We study Korean 4-ball billiards as a continuous, deterministic, sparse-reward RL problem with a fast exact simulator. Plain off-policy RL outperforms PPO, and random-start continue-on-miss training generalizes best, but the bare task still plateaus below one point per inning. Structure-free aids such as an action curriculum and learned reward-model bonus do not break this ceiling. Explicit geometry does: a first-contact aim constraint plus four carom features lifts SAC from 0.487 to 6.460 points/inning. The final jump comes from inference-time verification rather than reward design: greedy depth-2 lookahead uses the simulator to turn a ~ 6 -point policy into chains up to 8,451 consecutive scoring shots at 99.9% per-shot success. In this domain, learning supplies the prior, geometry makes it competent, and simulator search scales it.

1 The Korean 4-ball domain and learning environment

Game. Korean 4-ball (*sagu*) is a pocketless carom game played on a rectangular table with four balls: two cue balls (a white *player* cue and a yellow *opponent* cue) and two red object balls. A player scores exactly **one point** when their cue ball, struck once, caroms off **both** red balls in a single shot; contacting the opponent’s cue ball is a **foul**. After a scoring shot the player continues from the resulting table state, so play is a streak of shots terminated by the first miss or foul. The game is a pure single-agent control problem—there is no adversarial interaction through the table state—which makes it a clean testbed for continuous-control RL with a physically grounded, sparse objective.

Environment. We implement a NumPy physics simulator and expose it as a Gym environment (Billiards4BallInningEnv). The simulator is a from-scratch reimplement of standard billiards-physics models: the cue-tip impact follows Marlow’s level-cue instantaneous-point model [1], free flight uses the classic rigid-sphere slip/roll friction transition (reproduced in [1,7]), ball–cushion rebound uses the Han carom model [2], and ball–ball contact is the standard equal-mass restitution exchange; coefficients and the closed forms follow the open-source PoolTool simulator [6], which we port and strip down to the four modules 4-ball play requires (no pockets, jumps, or cloth effects) for a $\sim 20\times$ speedup. The **observation** is a 28-dimensional vector (4 balls $\times$ 7 state features: planar position, linear velocity, and angular velocity). The **action** is a 4-dimensional continuous control (θ, p, a, b) : aim angle $\theta \in [0, 2\pi)$, relative power $p \in [0, 1]$, and side/vertical cue-tip offsets $a, b \in [-1, 1]$ producing spin (projected to the unit disk to forbid miscues). The **reward** is the rule-defined carom score $\{0, 1\}$ per shot; no shaping is used unless stated. An episode is an *inning* of up to `max_shots`= 10 shots; with *continue-on-miss* the inning always runs the full budget, exposing the policy to a diverse stream of mid-rack states. One environment step equals one shot, so a budget of, e.g., 10^6 steps means one million simulated shots. Unless noted we evaluate the **mean points per inning** over 100 innings with `max_shots`= 10, and additionally report per-shot foul and success

rates. One shot costs ≈ 5 ms on a single CPU core, so 2×10^5 transitions take ≈ 15 –20 minutes on a laptop.

2 Algorithms and learning paradigms on the sparse task

2.1 SAC vs. PPO vs. TD3

We first compare three standard algorithms on the *plain* environment (random start, continue-on-miss, no domain knowledge, reward = raw $\{0, 1\}$ score). All are trained to 400k steps—the learning curves are already in their knee region by then, so the budget does not change the ranking—over 5 seeds.

Table 1: Algorithm comparison on the plain sparse task (5 seeds, 400k steps, eval = 100 innings, max_shots= 10). *mean*= points/inning; *max*= best single inning over all seeds. **Take-away: off-policy methods beat PPO by $\sim 2.5\times$ on score; SAC and TD3 are statistically indistinguishable but SAC is more stable.**

algo	mean $\pm$ std	max	foul/inn (%)
TD3	0.460$\pm$0.114	4	84
SAC	0.418$\pm$0.038	4	90
PPO	0.170$\pm$0.040	3	70

Off-policy SAC/TD3 reach 0.42–0.46 mean points, about $2.5\times$ PPO’s 0.17, and only the off-policy methods ever reach a 4-point inning. PPO plateaus by 200k (curve flat at ~ 0.15) while the off-policy pair keeps climbing in a separate band, so no ranking reversal is possible at larger budgets: every transition PPO collects is used once and discarded, which is fatal when scoring transitions are rare, whereas SAC/TD3 replay them. TD3’s mean is marginally higher but its seed variance is large (per-seed 0.33–0.62); SAC is markedly more stable (± 0.038). All methods stay weak overall—with foul_penalty= 0 fouls are unpenalized, so 70–90% of innings contain at least one foul and the best innings remain short—exposing the ceiling of the plain sparse task and motivating Sections 3–4.

2.2 Four training paradigms: start $\times$ miss handling

The environment admits four training regimes from the cross of *canonical* vs. *random* starting rack and *reset-on-miss* vs. *continue-on-miss*. We train SAC (400k, 3 seeds) under each and re-evaluate *all four under one unified protocol* (continue-on-miss, max_shots= 10, 100 innings) so they are comparable. The random-start eval measures generalization and is the primary metric.

Table 2: Training-paradigm comparison (SAC, 3 seeds, unified eval). **Take-away: random-start + continue-on-miss generalizes best and most stably; continue-on-miss is the dominant factor; canonical+reset overfits and collapses off its training rack.**

method	random-start mean $\pm$ std	canonical mean	foul%/shot	success%/shot
canonical, reset	0.073 $\pm$ 0.045	1.00	9.3	0.7
canonical, continue	0.353 $\pm$ 0.101	1.33	20.4	3.5
random, reset	0.407 $\pm$ 0.046	0.67	18.6	4.1
random, continue	0.453$\pm$0.012	0.67	20.3	4.5

Random + continue wins on the generalization metric (0.453) and is the most stable (± 0.012). Continue-on-miss is the dominant lever: switching reset $\rightarrow$ continue helps under both start regimes by exposing many mid-rack states per episode. Canonical+reset overfits its single rack and collapses to 0.073 on random racks. A key caution emerges from the training curves: ranked by trainer-reported episode return, *canonical+continue* (1.44) appears to dominate *random+continue* (0.37), but the unified eval *reverses* this—training-protocol return is a misleading model-selection signal across paradigms. We adopt random-start + continue-on-miss for all subsequent experiments.

3 Baseline and two attempts to improve it

We fix **SAC** (plain environment, random-start + continue-on-miss, 1M steps, 3 seeds) as the baseline because it is the most stable strong algorithm in Section 2. More budget does *not* break the ceiling: the 3-seed training curve knees near $\sim 250k$ and stays in a 0.40–0.47 band, while deterministic evaluation reaches only 0.487 ± 0.097 points/inning. The bottleneck is missing structure, not training time, so we test two structure-free ways to help the baseline.

3.1 Attempt 1: staged action curriculum (failed)

We first release the 4-D action space in stages: aim only, aim plus power, then the full action (θ, p, a, b) , warm-starting the same SAC policy across stages for 1M total steps. This does not help.

Table 3: Curriculum and RM-guided SAC vs. the plain SAC baseline (1M steps, 3 seeds). **Take-away: the curriculum underperforms the baseline; the RM lifts the mean somewhat but at a large extra data cost.**

method	mean $\pm$ std	max	rand. eval	extra cost
plain baseline	0.487 $\pm$ 0.097	4	0.053	—
staged curriculum	0.407 $\pm$ 0.052	4	0.053	none
RM-guided	0.840$\pm$0.000	4	0.087	10^7 offline shots

The curriculum is a clear *failure*: mean falls to 0.407 ± 0.052 , below the plain baseline, and the random-start eval stays at 0.053. The fixed release order spends budget on constrained sub-problems whose optima are not useful warm-starts for the full action space.

3.2 Attempt 2: learned reward-model (RM) guidance

We also train a small state–action model $V_{rm}(s, a) \approx P(\text{score} = 1 \mid s, a)$ from 10^7 random offline shots and add it as a dense SAC bonus; evaluation still uses the true binary carom score. The bonus raises train mean to 0.840, but the best inning remains 4 and random-start eval is only 0.087. More importantly, the gain is expensive and does not compound with real geometry: RM plus aim constraint reaches 2.357 ± 0.139 , below the constraint alone at 2.570. We therefore treat RM guidance as another negative result and turn to explicit domain knowledge.

4 Adding domain knowledge

Sections 2–3 establish that neither more budget, a curriculum, nor a learned reward model breaks the plain-RL ceiling. We now add *explicit* domain knowledge—a geometric action constraint and a few hand-chosen observation features—to the same SAC baseline (1M steps, 3 seeds, random-start + continue-on-miss). Table 4 reports the cumulative effect.

Table 4: Domain knowledge on the SAC baseline (1M steps, 3 seeds, training-protocol eval). *foul/sh* and *succ/sh* are per-shot foul and scoring probabilities. **Take-away: the aim constraint and the extra geometric features each lift the entire score distribution—together they take the agent from 0.49 to 6.46 points/inning.**

variant	mean $\pm$ std	max	foul/sh (%)	succ/sh (%)
plain SAC	0.487 $\pm$ 0.097	4	20.5	4.9
+ aim constraint	2.570 $\pm$ 0.128	6	15.4	25.7
+ extra feat.	6.460$\pm$0.123	10	11.1	64.6

4.1 Aim constraint

The cue’s aim θ is the one action dimension where most of the range is wasted: from a random layout the vast majority of directions send the cue ball into empty space, never contacting a red.

We therefore reparameterise θ as an *offset* within the angular window target $\pm \arcsin(2r/d)$ that geometrically *guarantees* first contact with the nearest red ball (radius r , distance d); the policy now chooses *where within* the guaranteed-contact cone to aim rather than whether to make contact at all. This single change is decisive: mean rises $0.487 \rightarrow 2.570$ ($5.3\times$), per-shot scoring rises from 4.9% to 25.7%, and—because most shots now at least touch a red instead of fouling—the per-shot foul rate *drops* ($20.5\% \rightarrow 15.4\%$) even though no penalty was added.

4.2 Extra geometric features

On top of the constraint we append four hand-chosen features to the observation, all measured from the red ball nearest the cue: the distances d_1, d_2 from the cue to the two red balls, and $\sin \varphi, \cos \varphi$ of the angle φ between the cue-to-nearest-red direction and the nearest-to-other-red direction. These encode the carom geometry the policy would otherwise have to infer from raw coordinates. The effect is the largest single lever in the paper (Table 4, three-way plain $\rightarrow$ constraint $\rightarrow$ constraint+extra): mean $2.570 \rightarrow 6.460$ ($2.5\times$), per-shot scoring $25.7\% \rightarrow 64.6\%$, and the best inning hits the 10-shot cap. Making the relevant geometry explicit matters far more than any change to the reward.

4.3 Reward shaping: foul penalty is the only useful knob

On the strong *constraint+extra* base we test three reward modifications (Table 5): a *gentle-shot* bonus for low-power scoring shots, a *near-miss* bonus for nearly completing the carom, and a fixed *foul penalty* of -1 . Neither dense-shaping term helps—gentle (6.407) and near-miss (6.353) are within noise of the 6.460 base, since a policy already scoring on 65% of shots has little headroom for an intermediate signal. The foul penalty is different: it lowers the mean ($6.460 \rightarrow 5.553$) but nearly *halves* the per-shot foul rate ($11.1\% \rightarrow 6.0\%$). It is thus a genuine score-vs-safety knob, and the one shaping term worth keeping: when fouls are costly (e.g. match play, where a foul cedes the turn) the foul penalty buys a large reduction in fouls for a modest scoring loss.

Table 5: Reward shaping on the constraint+extra base (1M steps, 3 seeds). **Take-away: dense shaping (gentle, near-miss) does not beat the base; only the foul penalty matters—it trades ~ 0.9 points/inning for a $\sim 45\%$ cut in the per-shot foul rate.**

variant	mean $\pm$ std	foul/sh (%)	succ/sh (%)
constraint+extra (base)	6.460$\pm$0.123	11.1	64.6
+ gentle shot	6.407 $\pm$ 0.164	9.9	64.1
+ near-miss	6.353 $\pm$ 0.100	10.2	63.5
+ foul penalty	5.553 $\pm$ 0.162	6.0	55.5

5 Inference-time lookahead search

The trained policy outputs a Gaussian over actions, but because our simulator runs a shot in ≈ 5 ms we need not commit to the policy mean: we can sample several candidate actions, simulate each, and keep the one with the best simulated outcome. This is the AlphaZero recipe—learned prior plus inference-time search—specialised to a continuous action space with a known, fast forward model. We use it to chase “infinite billiards”: starting from a random rack and terminating on the *first* miss or foul, how long a scoring chain can one inning sustain? The metric is the chain length (scoring shots before the first failure); we cap innings at 10,000 shots, a limit that is rarely reached.

A chain-friendly base policy. Long chains need *position control*: each shot should not only score but leave the cue ball where the next shot is easy. We therefore base the search on a 2M ensemble (3 seeds) trained with `constrain_aim + extra_features` plus two position-shaping terms—a *gentle-shot* bonus and a *setup* bonus that rewards finishing near the next red—and a mild `foul_penalty= 0.2` (rather than 1.0, which makes the policy too risk-averse to attempt chains). This nuances Section 4: shaping looked redundant on the 10-shot inning eval, but the long-chain regime is exactly where leave-position shaping pays off. The three seeds propose candidates jointly (mixing seeds covers more scoring modes at a fixed budget), and the *simulator itself is the verifier*—no

learned reward model is in the loop (the extra-feature policy uses a 32-d observation, incompatible with our 28-d $V(s, a)$ RM).

Search. Our search is a greedy depth-2 K -expansion: at state s_t the three policies emit K_1 candidates jointly, we simulate all from a snapshot and keep the top M_1 by immediate reward, expand each with K_2 children, simulate M_2 , and execute $\arg \max_{a_1} (r_1 + \gamma \max r_2)$ with $\gamma=0.99$. Because the environment is deterministic a candidate is fully described by its action—a second visit reproduces an identical trajectory—so there is no value in the visit-count revisiting or tree reuse that PUCT [5] relies on; freshly sampling K candidates at every shot uses the simulator budget far more effectively. The only knob that matters is therefore how many candidates we can afford per shot.

Results: chain length scales with inference compute. Lookahead transforms the agent. The same policy that scores ≈ 6 points per 10-shot inning unaided (Section 4) sustains chains of *thousands* of consecutive scoring shots, and the chain length grows monotonically with the per-shot simulation budget (Table 6). At the largest budget (400 simulations/shot) the best inning reaches an **8451-shot chain**— $4.2\times$ the prior multi-seed record (2000) and $11\times$ the best single-policy run (742) on this task. Per-shot success climbs to 99.9%: once in a chain the agent essentially never fails, so the failures that end an inning are almost always an unlucky *opening* rack (the zero-length episodes), which is why the mean chain sits well below the max. The trend is the headline: spending more compute at inference time—not a cleverer reward—monotonically buys longer chains.

Table 6: Greedy depth-2 lookahead on the 2M chain ensemble, swept over per-shot simulation budget (random-start, terminate on first miss/foul, max_shots= 10,000, 10 episodes each; chain length in scoring shots). **Take-away: chain length scales monotonically with inference-time compute, reaching an 8,451-shot chain at 99.9% per-shot success.**

K_1 total	sims/shot	mean chain	max chain	succ/shot (%)	wall/ep (min)
30	45	155	409	99.4	0.3
60	90	395	1771	99.7	1.7
150	225	1235	5716	99.9	14.2
300	400	1584	8451	99.9	28.0

6 Conclusion: a bitter lesson

Korean 4-ball is an unusual RL target. Unlike the discrete games where search was born or the dense-reward, model-free benchmarks where learning shines alone, it is *continuous, deterministic, and sparsely rewarded*, and it comes with a cheap, exact forward model. We set out to solve it with reinforcement learning, and the arc of our attempts reads as a small re-enactment of Sutton’s bitter lesson.

The hand-built, knowledge-injecting ideas disappointed. Plain RL plateaued below 0.5 points/inning regardless of budget (Section 3); a staged action curriculum did *worse* than no curriculum; and a learned reward model—precisely the kind of human-designed scaffolding the bitter lesson warns against—cost an extra 10^7 simulated shots to build yet added little and vanished entirely once a real lever was present. Even the explicit domain knowledge that *did* help (the aim constraint and geometric features, which lifted the agent from 0.5 to 6.5 points/inning) behaves like classic hand-engineering: a large but *bounded* one-time gain (Section 4).

What actually broke the ceiling were the two general methods that scale with computation. Learning gave us a competent prior; *search* did the rest. A simple greedy depth-2 lookahead—no learned value function, no tree, just the fast simulator used as its own verifier—turned a ~ 6 -point policy into chains of thousands of consecutive scoring shots (up to 8,451 at 99.9% per-shot success), with chain length scaling monotonically as we spent more compute per shot (Section 5). The lesson we draw is domain-specific but pointed: when an exact, fast forward model is available, the highest-leverage move is not a cleverer reward but spending compute at inference time. We tried to make 4-ball an RL problem; in the end, search heuristics solved it.

References

- [1] Marlow, W.C. (1995). *The Physics of Pocket Billiards*. MAST Publications.
- [2] Han, I. (2005). Dynamics in carom and three cushion billiards. *J. Mechanical Science and Technology*, 19(4):976–984.
- [3] Haarnoja, T., Zhou, A., Abbeel, P., Levine, S. (2018). Soft actor-critic. *ICML*.
- [4] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., Klimov, O. (2017). Proximal policy optimization. arXiv:1707.06347.
- [5] Silver, D., Hubert, T., Schrittwieser, J., et al. (2017). Mastering chess and shogi by self-play with a general RL algorithm. arXiv:1712.01815.
- [6] Kiefl, E. (2024). Pooltool. *JOSS*, 9(101):7301.
- [7] Alciatore, D.G. (2004). *The Illustrated Principles of Pool and Billiards*. Sterling Publishing.